

Applied Statistical Methods

Introduction to Python - Part IV

Xuemao Zhang
East Stroudsburg University

January 30, 2023

Outline

- Relative path
- Pandas series
- Data Frame
- Data Input
- Data output
- Summarizing data
- Sub-setting a data frame

Relative path

- 'Reading in' data is the first step of any data project/analysis

Relative vs. absolute paths (From Wiki)

- An **absolute or full path** points to the same location in a file system, regardless of the current working directory. To do that, it must include the **root directory**.
- This means if I try your code, and you use absolute paths, it won't work unless we have the exact same folder structure where Python is looking (bad).
- By contrast, a **relative path starts from some given working directory**, avoiding the need to provide the full absolute path. A filename can be considered as a relative path based at the current working directory.

Relative path

- Using relative paths simplifies the code since you don't need to write the full absolute path in order to find a file or directory in your Python project.
- Typical directory structure syntax applies
 - ▶ “../” - goes up one level
 - ▶ “./” - is the current directory
 - ▶ “~” - is your “home” directory

pandas series

- pandas have two data structures
 - ▶ Series
 - ▶ DataFrame
- A **Series** is a one-dimensional array-like object containing an array of data and an associated array of **data labels**, called its **index**.

```
import pandas as pd
obj = pd.Series([4, 7, -5, 3])
print(obj)
```

```
## 0    4
## 1    7
## 2   -5
## 3    3
## dtype: int64
```

pandas series

- `pandas.Series.values` return Series as ndarray or ndarray-like depending on the dtype.
 - ▶ <https://pandas.pydata.org/docs/reference/api/pandas.Series.values.html>

```
obj.values #it is a numpy.ndarray
```

```
## array([ 4,  7, -5,  3], dtype=int64)
```

```
obj.index
```

```
## RangeIndex(start=0, stop=4, step=1)
```

pandas series

- A Series can be created with a dictionary

```
sdata = {'Ohio': 35000, 'Texas': 71000, 'Oregon': 16000, 'Utah': 5000}  
obj2 = pd.Series(sdata)  
obj2.values
```

```
## array([35000, 71000, 16000,  5000], dtype=int64)
```

```
obj2.index
```

```
## Index(['Ohio', 'Texas', 'Oregon', 'Utah'], dtype='object')
```

Data Frame

- Data scientist are more comfortable dealing with tidy (Wickham and Grolemund, [R for Data Science](#), 2017) data frames:
 - ▶ Each variable must have its own column.
 - ▶ Each observation must have its own row.
 - ▶ Each value must have its own cell.

country	year	cases	population
Afghanistan	1999	18215	15467071
Afghanistan	2000	18666	20085360
Brazil	1999	31737	17206362
Brazil	2000	84488	17404898
China	1999	212258	1272015272
China	2000	213366	128553583

variables

country	year	cases	population
Afghanistan	1999	18215	15467071
Afghanistan	2000	18666	20085360
Brazil	1999	31737	17206362
Brazil	2000	84488	17404898
China	1999	212258	1272015272
China	2000	213366	128553583

observations

country	year	cases	population
Afghanistan	1999	18215	15467071
Afghanistan	2000	18666	20085360
Brazil	1999	31737	17206362
Brazil	2000	84488	17404898
China	1999	212258	1272015272
China	2000	213366	128553583

values

Data Frame

- One way to create a pandas DataFrame is convert dictionaries

```
df=pd.DataFrame({'col1':[1,3,11,2], 'col2':[9,13,24,16]})  
print(type(df))
```

```
## <class 'pandas.core.frame.DataFrame'>
```

```
print(df)
```

```
##      col1  col2  
## 0         1     9  
## 1         3    13  
## 2        11    24  
## 3         2    16
```

Data Frame

- Again, `pandas.DataFrame.values` return a *Numpy* representation of the DataFrame
 - ▶ Only the values in the DataFrame will be returned, the axes labels will be removed.
 - ▶ <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.values.html>

```
df.index
```

```
## RangeIndex(start=0, stop=4, step=1)
```

```
df.values # a numpy.ndarray
```

```
## array([[ 1,  9],  
##        [ 3, 13],  
##        [11, 24],  
##        [ 2, 16]], dtype=int64)
```

Data Frame

- Or the method `pandas.DataFrame.to_numpy` can be used
 - ▶ https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.to_numpy.html

```
df2=df.to_numpy()  
print(df2)
```

```
## [[ 1  9]  
##   [ 3 13]  
##  [11 24]  
##   [ 2 16]]
```

Data Input

- The library `pandas` provides powerful tools for data import.
- We are going to focus on simple delimited files first
 - ▶ comma separated (e.g. `'csv'`)
 - ▶ tab delimited (e.g. `'txt'`)
 - ▶ Microsoft excel (e.g. `'xlsx'`)
- Throughout this course, we will be working with CSV files.

Data Input

- Consider the data set `earthquakes.csv` from https://raw.githubusercontent.com/stefmolin/Hands-On-Data-Analysis-with-Pandas-2nd-edition/master/ch_02/data/earthquakes.csv which is also downloaded to the folder `data`.
- https://pandas.pydata.org/docs/reference/api/pandas.read_csv.html
- The imported file is a `pandas.DataFrame`

```
earthquake= pd.read_csv("../data/earthquakes.csv", header=0)  
# the data can be read from the website directly  
print(type(earthquake))
```

```
## <class 'pandas.core.frame.DataFrame'>
```

Data Input

- column names.

```
vars=earthquake.columns  
print(vars)
```

```
## Index(['alert', 'cdi', 'code', 'detail', 'dmin', 'felt', 'gap', 'ids', '  
##      'magType', 'mmi', 'net', 'nst', 'place', 'rms', 'sig', 'sources',  
##      'status', 'time', 'title', 'tsunami', 'type', 'types', 'tz', 'upd  
##      'url'],  
##      dtype='object')
```

```
print(type(vars))
```

```
## <class 'pandas.core.indexes.base.Index'>
```

Data Input

- `Index.to_list()` returns a list of the values.
 - ▶ https://pandas.pydata.org/docs/reference/api/pandas.Index.to_list.html

```
vars=earthquake.columns.to_list()
print(type(vars))
```

```
## <class 'list'>
```

```
print(vars)
```

```
## ['alert', 'cdi', 'code', 'detail', 'dmin', 'felt', 'gap', 'ids', 'mag',
```

Data Input

- Data types of the columns

```
print(earthquake.dtypes)
```

```
## alert      object
## cdi        float64
## code       object
## detail     object
## dmin       float64
## felt       float64
## gap        float64
## ids        object
## mag        float64
## magType    object
## mmi        float64
## net        object
## nst        float64
## place      object
## rms        float64
## sig        int64
## sources    object
## status     object
```


Data Input

- The following are some commonly used data frame attributes

Attribute	Returns
<code>dtypes</code>	The data types of each column
<code>shape</code>	Dimensions of the <code>DataFrame</code> object in a tuple of the form (number of rows, number of columns)
<code>index</code>	The <code>Index</code> object along the rows of the <code>DataFrame</code> object
<code>columns</code>	The name of the columns (as an <code>Index</code> object)
<code>values</code>	The data in the <code>DataFrame</code> object
<code>empty</code>	Check if the <code>DataFrame</code> object is empty

Data Input

- Check if the data frame is empty

```
print(earthquake.empty)
```

```
## False
```

- The data in the data frame

```
print(earthquake.values) #a 'numpy.ndarray'
```

```
## [[nan nan '37389218' ... -480.0 1539475395144
##    'https://earthquake.usgs.gov/earthquakes/eventpage/ci37389218']
## [nan nan '37389202' ... -480.0 1539475253925
##    'https://earthquake.usgs.gov/earthquakes/eventpage/ci37389202']
## [nan 4.4 '37389194' ... -480.0 1539536756176
##    'https://earthquake.usgs.gov/earthquakes/eventpage/ci37389194']
## ...
## [nan nan '2018261000' ... -240.0 1537243777410
##    'https://earthquake.usgs.gov/earthquakes/eventpage/pr2018261000']
## [nan nan '38063959' ... -480.0 1537230211640
##    'https://earthquake.usgs.gov/earthquakes/eventpage/ci38063959']
## [nan nan '38063935' ... -480.0 1537305830770
##    'https://earthquake.usgs.gov/earthquakes/eventpage/ci38063935']]
```

Data Input

- Check the dimensions

- ▶ <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.shape.html>

```
dim=earthquake.shape  
print(type(dim))
```

```
## <class 'tuple'>
```

```
print(dim)
```

```
## (9332, 26)
```

Data Input

- Check the row names/index

- ▶ <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.index.html#>

```
rownames=earthquake.index  
print(rownames)
```

```
## RangeIndex(start=0, stop=9332, step=1)
```

Data Input

- Get the head using the head() method
- <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.head.html>

```
print(earthquake.head()) #Return the first n rows
```

```
##      alert  cdi  ...      updated
## 0      NaN  NaN  ...  1539475395144  https://earthquake.usgs.gov/earthquake
## 1      NaN  NaN  ...  1539475253925  https://earthquake.usgs.gov/earthquake
## 2      NaN  4.4  ...  1539536756176  https://earthquake.usgs.gov/earthquake
## 3      NaN  NaN  ...  1539475196167  https://earthquake.usgs.gov/earthquake
## 4      NaN  NaN  ...  1539477547926  https://earthquake.usgs.gov/earthquake
##
## [5 rows x 26 columns]
```

Data Input

- Get the tail using the `tail()` method
- <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.tail.html>

```
print(earthquake.tail()) #Return the last n rows
```

```
##          alert    ...                               url
## 9327      NaN    ...  https://earthquake.usgs.gov/earthquakes/eventp...
## 9328      NaN    ...  https://earthquake.usgs.gov/earthquakes/eventp...
## 9329      NaN    ...  https://earthquake.usgs.gov/earthquakes/eventp...
## 9330      NaN    ...  https://earthquake.usgs.gov/earthquakes/eventp...
## 9331      NaN    ...  https://earthquake.usgs.gov/earthquakes/eventp...
##
## [5 rows x 26 columns]
```

Data Input

- We can use the `info()` method to see how many non-null entries of each column we have and get information on our index.

```
print(earthquake.info()) #Return the first n rows
```

```
## <class 'pandas.core.frame.DataFrame'>
## RangeIndex: 9332 entries, 0 to 9331
## Data columns (total 26 columns):
## #      Column      Non-Null Count  Dtype
## ---  -
## 0      alert        59 non-null    object
## 1      cdi            329 non-null   float64
## 2      code           9332 non-null  object
## 3      detail        9332 non-null  object
## 4      dmin          6139 non-null  float64
## 5      felt          329 non-null   float64
## 6      gap           6164 non-null  float64
## 7      ids           9332 non-null  object
## 8      mag           9331 non-null  float64
## 9      magType       9331 non-null  object
## 10     mmi           93 non-null    float64
## 11     net           9332 non-null  object
```

Data Input

- `read_csv` can be used to read a `.txt` file
 - ▶ `quiz1.txt` is a *tab-delimited* file

```
quiz=pd.read_csv("../data/quiz1.txt",header=0, delimiter='\t')
print(quiz.head())
```

##	ID	Q1	Q2	Q3	Q4	Q5	Q6
## 0	1	8.0	9.0	10.0	9.5	10.0	8.0
## 1	2	8.0	8.0	8.0	10.0	9.0	8.0
## 2	3	10.0	7.0	10.0	10.0	10.0	8.0
## 3	4	6.0	5.0	9.0	8.0	5.0	NaN
## 4	5	10.0	6.0	8.0	6.0	NaN	NaN

Data Input

- `read_excel` is used to read in Excel files
- https://pandas.pydata.org/docs/reference/api/pandas.read_excel.html?highlight=read_excel
 - ▶ Install `xlrd >= 1.0.0` for Excel support Use `pip` or `conda` to install `xlrd`.

```
quiz2=pd.read_excel("../data/quiz2.xls",sheet_name=0, header=0)
print(quiz2.head())
```

##	ID	Q1	Q2	Q3	Q4	Q5	Q6
## 0	1	8.0	9.0	10.0	9.5	10.0	8.0
## 1	2	8.0	8.0	8.0	10.0	9.0	8.0
## 2	3	10.0	7.0	10.0	10.0	10.0	8.0
## 3	4	6.0	5.0	9.0	8.0	5.0	NaN
## 4	5	10.0	6.0	8.0	6.0	NaN	NaN

Data output

- Sometimes we want to save our data frame to a file so that we can share it with others.

- ▶ [https:](https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.to_csv.html)

- [//pandas.pydata.org/docs/reference/api/pandas.DataFrame.to_csv.html](https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.to_csv.html)

```
df.to_csv('output.csv', index=False) #remove the row labels
```

- The following code save the data quiz to the file quiz2.csv

```
quiz.to_csv('quiz3.csv', sep=',', header=True)
```

- Check out the following resource in the pandas documentation for the full list of capabilities:

- https://pandas.pydata.org/pandas-docs/stable/user_guide/io.html

Summarizing data

- Method `describe()` gives us the 5-number summary, along with the count, mean, and standard deviation of the numeric columns

```
import pandas as pd
iris=pd.read_csv("../data/iris.csv",header=0, delimiter=',')
print(iris.describe())
```

	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width
## count	150.000000	150.000000	150.000000	150.000000
## mean	5.843333	3.057333	3.758000	1.199333
## std	0.828066	0.435866	1.765298	0.762238
## min	4.300000	2.000000	1.000000	0.100000
## 25%	5.100000	2.800000	1.600000	0.300000
## 50%	5.800000	3.000000	4.350000	1.300000
## 75%	6.400000	3.300000	5.100000	1.800000
## max	7.900000	4.400000	6.900000	2.500000

Summarizing data

- By default, `describe()` won't give us any information about the columns of type object, but we can either provide `include='all'` as an argument or run it separately for the data of type `np.object`

```
print(iris.describe(include='all'))
```

##	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
## count	150.000000	150.000000	150.000000	150.000000	150
## unique	NaN	NaN	NaN	NaN	3
## top	NaN	NaN	NaN	NaN	setosa
## freq	NaN	NaN	NaN	NaN	50
## mean	5.843333	3.057333	3.758000	1.199333	NaN
## std	0.828066	0.435866	1.765298	0.762238	NaN
## min	4.300000	2.000000	1.000000	0.100000	NaN
## 25%	5.100000	2.800000	1.600000	0.300000	NaN
## 50%	5.800000	3.000000	4.350000	1.300000	NaN
## 75%	6.400000	3.300000	5.100000	1.800000	NaN
## max	7.900000	4.400000	6.900000	2.500000	NaN

Summarizing data

- By default, `describe()` won't give us any information about the columns of type object, but we can either provide `include='all'` as an argument or run it separately for the data of type object

```
print(iris.describe(include=object))
```

```
##           Species
## count         150
## unique          3
## top      setosa
## freq         50
```

Summarizing data

- Sometimes, we just want a particular *statistic* for a specific column(s)
- A list of *calculation methods* for data frames

Method	Description	Data types
<code>count()</code>	The number of non-null observations	Any
<code>nunique()</code>	The number of unique values	Any
<code>sum()</code>	The total of the values	Numerical or Boolean
<code>mean()</code>	The average of the values	Numerical or Boolean
<code>median()</code>	The median of the values	Numerical
<code>min()</code>	The minimum of the values	Numerical
<code>idxmin()</code>	The index where the minimum values occurs	Numerical
<code>max()</code>	The maximum of the values	Numerical
<code>idxmax()</code>	The index where the maximum value occurs	Numerical
<code>abs()</code>	The absolute values of the data	Numerical
<code>std()</code>	The standard deviation	Numerical
<code>var()</code>	The variance	Numerical
<code>cov()</code>	The covariance between two <code>Series</code> , or a covariance matrix for all column combinations in a <code>DataFrame</code>	Numerical
<code>corr()</code>	The correlation between two <code>Series</code> , or a correlation matrix for all column combinations in a <code>DataFrame</code>	Numerical
<code>quantile()</code>	Calculates a specific quantile	Numerical
<code>cumsum()</code>	The cumulative sum	Numerical or Boolean

Summarizing data

- Distinct values of a column

```
iris['Species'].nunique()  # iris.Species.nunique()
```

```
## 3
```

```
iris['Species'].unique()
```

```
## array(['setosa', 'versicolor', 'virginica'], dtype=object)
```

- Frequency table

```
iris['Species'].value_counts()
```

```
## setosa      50
```

```
## versicolor  50
```

```
## virginica   50
```

```
## Name: Species, dtype: int64
```

Sub-setting a data frame

- We can drop a column using method `DataFrame.drop()`
 - ▶ `axis=1` means dropping columns
 - ▶ <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.drop.html>

```
iris.drop(['Petal.Length', 'Petal.Width'], axis=1, inplace=True)  
iris.columns
```

```
## Index(['Sepal.Length', 'Sepal.Width', 'Species'], dtype='object')
```


Sub-setting a data frame

- Select columns
 - ▶ using indexing
 - ▶ using dot notation

```
ID=quiz['ID']  
print(type(ID))
```

```
## <class 'pandas.core.series.Series'>
```

```
print(ID)
```

```
## 0      1  
## 1      2  
## 2      3  
## 3      4  
## 4      5  
## 5      6  
## 6      7  
## 7      8  
## 8      9  
## 9     10  
## 10     11  
## 11     12
```

Sub-setting a data frame

- Select columns
 - ▶ using indexing
 - ▶ using dot notation

```
quiz_sub=quiz[["ID", "Q1", "Q2", "Q3"]]  
print(type(quiz_sub))
```

```
## <class 'pandas.core.frame.DataFrame'>
```

```
quiz_sub.head()
```

```
##      ID      Q1      Q2      Q3  
## 0      1      8.0      9.0     10.0  
## 1      2      8.0      8.0      8.0  
## 2      3     10.0      7.0     10.0  
## 3      4      6.0      5.0      9.0  
## 4      5     10.0      6.0      8.0
```

Sub-setting a data frame

- Select columns
 - ▶ using indexing
 - ▶ using dot notation (it is like the \$ in R)

```
ID2=quiz.ID  
print(type(ID2))
```

```
## <class 'pandas.core.series.Series'>
```

```
print(ID2)
```

```
## 0      1  
## 1      2  
## 2      3  
## 3      4  
## 4      5  
## 5      6  
## 6      7  
## 7      8  
## 8      9  
## 9     10  
## 10     11  
## 11     12
```

Sub-setting a data frame

- Subsetting rows (or called slicing)

```
quiz[10:20]
```

##		ID	Q1	Q2	Q3	Q4	Q5	Q6
##	10	11	10.0	NaN	4.0	7.0	9.0	NaN
##	11	12	6.0	8.0	5.0	8.0	10.0	8.0
##	12	13	10.0	10.0	8.0	10.0	10.0	9.0
##	13	14	10.0	10.0	10.0	10.0	9.0	8.0
##	14	15	10.0	10.0	10.0	10.0	10.0	10.0
##	15	16	10.0	8.0	9.0	6.0	9.0	6.0
##	16	17	10.0	NaN	9.0	NaN	6.0	9.0
##	17	18	10.0	10.0	10.0	7.0	6.0	9.0
##	18	19	10.0	7.0	7.0	7.0	10.0	8.0
##	19	20	10.0	10.0	10.0	10.0	10.0	8.0

Sub-setting a data frame

- Subsetting rows (or called slicing)
- `pandas.DataFrame.iloc` can be used to subset both rows and columns using **integer-based** lookups
 - ▶ <https://pandas.pydata.org/docs/reference/api/pandas.DataFrame.iloc.html>

```
import numpy as np
np.random.seed(1)
train = np.random.choice(40, 20, replace=False)
quiz.iloc[train, :]
```

##	ID	Q1	Q2	Q3	Q4	Q5	Q6
## 2	3	10.0	7.0	10.0	10.0	10.0	8.0
## 31	32	10.0	10.0	10.0	10.0	9.0	9.0
## 3	4	6.0	5.0	9.0	8.0	5.0	NaN
## 21	22	10.0	10.0	10.0	7.0	10.0	7.0
## 27	28	NaN	4.0	8.0	5.0	10.0	9.0
## 29	30	10.0	9.0	10.0	9.0	5.0	7.0
## 22	23	10.0	10.0	NaN	9.0	10.0	9.0
## 39	40	10.0	10.0	10.0	8.0	10.0	8.0
## 19	20	10.0	10.0	10.0	10.0	10.0	8.0
## 26	27	10.0	10.0	10.0	10.0	10.0	10.0
## 32	33	10.0	8.0	8.0	6.0	9.0	7.0

Sub-setting a data frame

- We can subset rows using logic values

```
select = np.in1d(range(40), train)
quiz[select]
```

##	ID	Q1	Q2	Q3	Q4	Q5	Q6
## 2	3	10.0	7.0	10.0	10.0	10.0	8.0
## 3	4	6.0	5.0	9.0	8.0	5.0	NaN
## 4	5	10.0	6.0	8.0	6.0	NaN	NaN
## 10	11	10.0	NaN	4.0	7.0	9.0	NaN
## 14	15	10.0	10.0	10.0	10.0	10.0	10.0
## 17	18	10.0	10.0	10.0	7.0	6.0	9.0
## 19	20	10.0	10.0	10.0	10.0	10.0	8.0
## 21	22	10.0	10.0	10.0	7.0	10.0	7.0
## 22	23	10.0	10.0	NaN	9.0	10.0	9.0
## 26	27	10.0	10.0	10.0	10.0	10.0	10.0
## 27	28	NaN	4.0	8.0	5.0	10.0	9.0
## 28	29	6.0	9.0	4.0	5.0	9.0	0.0
## 29	30	10.0	9.0	10.0	9.0	5.0	7.0
## 30	31	10.0	10.0	10.0	7.0	6.0	NaN
## 31	32	10.0	10.0	10.0	10.0	9.0	9.0
## 32	33	10.0	8.0	8.0	6.0	9.0	7.0

Sub-setting a data frame

- Check if two dataframes are equal

```
quiz.iloc[train, :].equals(quiz[select])
```

```
## False
```

- The two data frames will be identical if we sort the index
 - ▶ https://pandas.pydata.org/pandas-docs/stable/reference/api/pandas.DataFrame.sort_index.html

```
quiz_train=quiz.iloc[train, :].sort_index(axis = 0)  
quiz_train.equals(quiz[select])
```

```
## True
```

Sub-setting a data frame

```
quiz.iloc[range(10,15)]
```

##		ID	Q1	Q2	Q3	Q4	Q5	Q6
##	10	11	10.0	NaN	4.0	7.0	9.0	NaN
##	11	12	6.0	8.0	5.0	8.0	10.0	8.0
##	12	13	10.0	10.0	8.0	10.0	10.0	9.0
##	13	14	10.0	10.0	10.0	10.0	9.0	8.0
##	14	15	10.0	10.0	10.0	10.0	10.0	10.0

Sub-setting a data frame

- Subsetting discontinuous rows

```
quiz.iloc[ [*range(10,15),*range(20,25)],]
```

		ID	Q1	Q2	Q3	Q4	Q5	Q6
##	10	11	10.0	NaN	4.0	7.0	9.0	NaN
##	11	12	6.0	8.0	5.0	8.0	10.0	8.0
##	12	13	10.0	10.0	8.0	10.0	10.0	9.0
##	13	14	10.0	10.0	10.0	10.0	9.0	8.0
##	14	15	10.0	10.0	10.0	10.0	10.0	10.0
##	20	21	10.0	9.0	10.0	8.5	10.0	5.0
##	21	22	10.0	10.0	10.0	7.0	10.0	7.0
##	22	23	10.0	10.0	NaN	9.0	10.0	9.0
##	23	24	10.0	10.0	10.0	10.0	9.0	9.0
##	24	25	10.0	9.0	8.0	6.0	9.0	8.0

Sub-setting a data frame

- Or equivalently using index only

```
print(quiz.iloc[0:8, [0,1,3]])
```

##	ID	Q1	Q3
## 0	1	8.0	10.0
## 1	2	8.0	8.0
## 2	3	10.0	10.0
## 3	4	6.0	9.0
## 4	5	10.0	8.0
## 5	6	NaN	10.0
## 6	7	10.0	9.0
## 7	8	10.0	10.0

Sub-setting a data frame

- We can subset both rows and columns

```
quiz[["ID", "Q1", "Q2", "Q3"]][10:25]  
#quiz_sub=quiz[["ID", "Q1", "Q2", "Q3"]]  
#quiz_sub[10:25]
```

##		ID	Q1	Q2	Q3
##	10	11	10.0	NaN	4.0
##	11	12	6.0	8.0	5.0
##	12	13	10.0	10.0	8.0
##	13	14	10.0	10.0	10.0
##	14	15	10.0	10.0	10.0
##	15	16	10.0	8.0	9.0
##	16	17	10.0	NaN	9.0
##	17	18	10.0	10.0	10.0
##	18	19	10.0	7.0	7.0
##	19	20	10.0	10.0	10.0
##	20	21	10.0	9.0	10.0
##	21	22	10.0	10.0	10.0
##	22	23	10.0	10.0	NaN
##	23	24	10.0	10.0	10.0
##	24	25	10.0	9.0	8.0

Sub-setting a data frame

- Pandas indexing operations provide us with a one-method way to select both the rows and the columns we want.
- `pandas.DataFrame.loc` can be used to subset both rows and columns using **label-based** lookups.
- Using both labels and integer index using `loc`

```
print(quiz.loc[0:8, ["ID", "Q1", "Q3"]])
```

	ID	Q1	Q3
## 0	1	8.0	10.0
## 1	2	8.0	8.0
## 2	3	10.0	10.0
## 3	4	6.0	9.0
## 4	5	10.0	8.0
## 5	6	NaN	10.0
## 6	7	10.0	9.0
## 7	8	10.0	10.0
## 8	9	10.0	6.0

Sub-setting a data frame

- To look up cell values, we use `at[]` and `iat[]`, which are faster.

```
print(quiz.at[0, 'Q3'])
```

```
## 10.0
```

```
print(quiz.iat[0, 3])
```

```
## 10.0
```

Sub-setting a data frame

- Filtering: we can subset rows based on the restrictions or **Boolean masks** on a column(s)

```
quiz[quiz.Q6>= 9]
```

##	ID	Q1	Q2	Q3	Q4	Q5	Q6
## 8	9	10.0	10.0	6.0	7.0	8.0	10.0
## 12	13	10.0	10.0	8.0	10.0	10.0	9.0
## 14	15	10.0	10.0	10.0	10.0	10.0	10.0
## 16	17	10.0	NaN	9.0	NaN	6.0	9.0
## 17	18	10.0	10.0	10.0	7.0	6.0	9.0
## 22	23	10.0	10.0	NaN	9.0	10.0	9.0
## 23	24	10.0	10.0	10.0	10.0	9.0	9.0
## 26	27	10.0	10.0	10.0	10.0	10.0	10.0
## 27	28	NaN	4.0	8.0	5.0	10.0	9.0
## 31	32	10.0	10.0	10.0	10.0	9.0	9.0
## 33	34	10.0	10.0	6.0	8.0	8.0	9.0
## 34	35	10.0	10.0	7.0	8.0	10.0	10.0

Sub-setting a data frame

- Filtering: we can subset rows based on the restrictions or **Boolean masks** on a column(s)

```
quiz.loc[quiz.Q6 >= 9]
```

##	ID	Q1	Q2	Q3	Q4	Q5	Q6
## 8	9	10.0	10.0	6.0	7.0	8.0	10.0
## 12	13	10.0	10.0	8.0	10.0	10.0	9.0
## 14	15	10.0	10.0	10.0	10.0	10.0	10.0
## 16	17	10.0	NaN	9.0	NaN	6.0	9.0
## 17	18	10.0	10.0	10.0	7.0	6.0	9.0
## 22	23	10.0	10.0	NaN	9.0	10.0	9.0
## 23	24	10.0	10.0	10.0	10.0	9.0	9.0
## 26	27	10.0	10.0	10.0	10.0	10.0	10.0
## 27	28	NaN	4.0	8.0	5.0	10.0	9.0
## 31	32	10.0	10.0	10.0	10.0	9.0	9.0
## 33	34	10.0	10.0	6.0	8.0	8.0	9.0
## 34	35	10.0	10.0	7.0	8.0	10.0	10.0

Sub-setting a data frame

```
quiz.loc[(quiz.Q5>= 9) | (quiz.Q6>= 9)]
```

```
#quiz[(quiz.Q5>= 9) | (quiz.Q6>= 9)]
```

##	ID	Q1	Q2	Q3	Q4	Q5	Q6
## 0	1	8.0	9.0	10.0	9.5	10.0	8.0
## 1	2	8.0	8.0	8.0	10.0	9.0	8.0
## 2	3	10.0	7.0	10.0	10.0	10.0	8.0
## 5	6	NaN	9.0	10.0	10.0	10.0	NaN
## 6	7	10.0	5.0	9.0	8.0	10.0	7.0
## 7	8	10.0	10.0	10.0	9.0	10.0	7.0
## 8	9	10.0	10.0	6.0	7.0	8.0	10.0
## 9	10	8.0	10.0	10.0	4.0	10.0	7.0
## 10	11	10.0	NaN	4.0	7.0	9.0	NaN
## 11	12	6.0	8.0	5.0	8.0	10.0	8.0
## 12	13	10.0	10.0	8.0	10.0	10.0	9.0
## 13	14	10.0	10.0	10.0	10.0	9.0	8.0
## 14	15	10.0	10.0	10.0	10.0	10.0	10.0
## 15	16	10.0	8.0	9.0	6.0	9.0	6.0
## 16	17	10.0	NaN	9.0	NaN	6.0	9.0
## 17	18	10.0	10.0	10.0	7.0	6.0	9.0
## 18	19	10.0	7.0	7.0	7.0	10.0	8.0
## 19	20	10.0	10.0	10.0	10.0	10.0	8.0

Sub-setting a data frame

```
quiz.loc[(quiz.Q5>= 9) & (quiz.Q6>= 9)]  
#quiz[(quiz.Q5>= 9) & (quiz.Q6>= 9)]
```

##	ID	Q1	Q2	Q3	Q4	Q5	Q6
## 12	13	10.0	10.0	8.0	10.0	10.0	9.0
## 14	15	10.0	10.0	10.0	10.0	10.0	10.0
## 22	23	10.0	10.0	NaN	9.0	10.0	9.0
## 23	24	10.0	10.0	10.0	10.0	9.0	9.0
## 26	27	10.0	10.0	10.0	10.0	10.0	10.0
## 27	28	NaN	4.0	8.0	5.0	10.0	9.0
## 31	32	10.0	10.0	10.0	10.0	9.0	9.0
## 34	35	10.0	10.0	7.0	8.0	10.0	10.0

Sub-setting a data frame

- **Important note:** When creating Boolean masks, we must use bitwise operators (`&`, `|`, `~`) instead of logical operators (`and`, `or`, `not`).

Sub-setting a data frame

```
quiz.loc[(quiz.Q5>= 9) & (quiz.Q6>= 9),['Q1','Q2','Q3','Q4']]
```

##		Q1	Q2	Q3	Q4
##	12	10.0	10.0	8.0	10.0
##	14	10.0	10.0	10.0	10.0
##	22	10.0	10.0	NaN	9.0
##	23	10.0	10.0	10.0	10.0
##	26	10.0	10.0	10.0	10.0
##	27	NaN	4.0	8.0	5.0
##	31	10.0	10.0	10.0	10.0
##	34	10.0	10.0	7.0	8.0

Sub-setting a data frame

- Two Boolean masks can be avoided using the `between()` method

```
quiz.loc[quiz.Q5. between(5,6)]
```

##	ID	Q1	Q2	Q3	Q4	Q5	Q6
## 3	4	6.0	5.0	9.0	8.0	5.0	NaN
## 16	17	10.0	NaN	9.0	NaN	6.0	9.0
## 17	18	10.0	10.0	10.0	7.0	6.0	9.0
## 29	30	10.0	9.0	10.0	9.0	5.0	7.0
## 30	31	10.0	10.0	10.0	7.0	6.0	NaN
## 35	36	8.0	7.0	5.0	8.0	5.0	0.0

Sub-setting a data frame

- The `isin()` method can be used to create a Boolean mask for values that match one of a list of values

```
earthquake= pd.read_csv("../data/earthquakes.csv", header=0)
earthquake1=earthquake.loc[earthquake.magType.isin(['mw', 'mwb']),
    ['alert', 'mag', 'magType', 'title', 'tsunami', 'type']]
earthquake1.head()
```

```
##          alert  mag  ... tsunami      type
## 995      NaN  3.35  ...         0  earthquake
## 1465   green  3.83  ...         0  earthquake
## 2414   green  3.83  ...         1  earthquake
## 4988   green  4.41  ...         1  earthquake
## 6307   green  5.80  ...         0  earthquake
##
## [5 rows x 6 columns]
```

Sub-setting a data frame

- We can use `idxmin()` and `idxmax()` for the indices of the minimum and maximum values of a variable for filtering.

```
earthquake.loc[[earthquake.mag.idxmin(), earthquake.mag.idxmax()],  
['alert', 'mag', 'magType', 'title', 'tsunami', 'type']]
```

```
##      alert  mag  ... tsunami      type  
## 2409   NaN -1.26  ...        0  earthquake  
## 5263   red  7.50  ...        1  earthquake  
##  
## [2 rows x 6 columns]
```

License



This work is licensed under a [Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-nc-sa/4.0/).